

MINDPILOT BUSINESS PLAN



OmniPilot

Personal AI Assistant

Voice-first, local-first, multi-interface AI companion
integrated with the Sidewall product portfolio

Phase 1: Strategy

Phase 2: Business Design

Phase 3: Product Design

Phase 4: Execution

Prepared for: Soumya Swain | Sidewall / DevPilot
Date: April 16, 2026 | **Framework:** MindPilot 4-Phase
Status: Research & Planning Complete

Table of Contents

P1 Phase 1: Strategy

- 1.1 Problem Statement
- 1.2 Target Customer
- 1.3 Existing Alternatives & Competitors
- 1.4 Competitive Feature Heatmap
- 1.5 Unfair Advantage
- 1.6 Market Size (TAM / SAM / SOM)
- 1.7 Business Model
- 1.8 Risks & Mitigation
- 1.9 6-Month Success Metrics

P2 Phase 2: Business Design

- 2.1 Revenue Model & Unit Economics
- 2.2 Cost Structure & Path to Profitability
- 2.3 Customer Journey
- 2.4 Business Model Canvas
- 2.5 Channels & Partnerships
- 2.6 Key Metrics

P3 Phase 3: Product Design

- 3.1 System Architecture
- 3.2 Data Flow Pipeline
- 3.3 Memory Architecture
- 3.4 Tech Stack & RAM Budget
- 3.5 MVP Feature Set
- 3.6 Product Portfolio Integration

P4 Phase 4: Execution

- 4.1 Build Roadmap
- 4.2 Sprint 1 (Week 1-2) Scope
- 4.3 Launch Checklist
- 4.4 Go-to-Market Plan

Strategy: The Why & Where

1.1 What Problem Are We Solving?

PROBLEM STATEMENT

Business owners and professionals juggle 5-10 tools daily — WhatsApp, invoicing, ad dashboards, trading screens, sprint boards. Critical information gets lost across conversations, meetings, and apps. There is no single assistant that:

- Listens to your conversations and **remembers everything**
- Works across **all your business tools** with voice commands
- Keeps your data **100% private** (no cloud dependency)
- Is **affordable for Indian SMBs** (not \$20+/month Western pricing)

1.2 Who Has This Problem?


SEGMENT	PROFILE	PAIN LEVEL	WILLINGNESS TO PAY
SMB Owners (Primary)	Run 1-3 businesses, 5-20 employees, India Tier 1-2 cities	Very High	Rs 500-2,000/mo
Freelancers/Consultants	Client-facing, 5-15 active projects, need memory	High	Rs 300-500/mo
Existing Sidewall Users	BizBot/AdPilot/QuickBillPro customers	Very High	Rs 3,000-8,000/mo add-on
Knowledge Workers	Engineers, lawyers, doctors — heavy meeting load	High	Rs 500-1,000/mo
Traders	Active market participants, need ambient awareness	Moderate	Rs 500-1,500/mo

1.3 How Are They Solving It Today?

CURRENT SOLUTION	WHAT IT DOES	WHAT'S MISSING
Omi AI (\$89 pendant + \$16/mo)	Wearable that records & transcribes	No product integrations, no India pricing, cloud-only
Limitless (\$99 + \$19/mo)	Meeting memory + pendant	No WhatsApp, no action execution, Western pricing
Notion AI (\$10/mo)	Workspace AI over documents	Not a personal companion, no voice, no ambient listening

CURRENT SOLUTION	WHAT IT DOES	WHAT'S MISSING
ChatGPT/Claude	General AI chat	Forgets everything. No context. No integrations
Manual notes + memory	Human brain + notebooks	Doesn't scale. Forgets. Can't search. Can't act

1.4 Competitive Feature Heatmap



COMPETITIVE GAPS TO EXPLOIT

- **No WhatsApp integration** — None of the top 5 have native WhatsApp. Massive for India (78% SMBs use WhatsApp)
- **No action execution** — All 5 observe and summarize. None actually DO things (send invoices, reply to customers, execute trades)
- **No India features** — GST compliance, Hindi voice, INR trading, Indian city services — untouched by global players
- **Privacy vs. proactivity tradeoff** — Khoj is private but reactive. Others are proactive but cloud-only. OmniPilot can be both


1.5 Why Will They Switch to Us?

UNFAIR ADVANTAGE

1. **7-product integration** — No competitor connects to your invoicing + ads + WhatsApp CRM + trading + sprint management. OmniPilot does all 7
2. **Local-first privacy** — Everything on your Mac/phone. Competitors send audio to cloud. OmniPilot: zero cloud dependency
3. **India-first pricing** — Rs 499/month vs \$19/month competitors. Built for Indian SMBs, not Silicon Valley
4. **Action, not just memory** — "Send invoice to Brand Tusker" actually sends it. "Reply on WhatsApp" actually replies. Not just a note-taker
5. **Existing customer base** — BizBot and AdPilot customers are day-1 prospects. Upsell, not cold acquisition

1.6 Market Size

AI Assistant Market Sizing



\$37.7B

AI COMPANION TAM

\$5.1B

INDIA AI SAM

\$1.12B

INDIA CPAAS SOM

3x

GROWTH BY 2027

1.7 Business Model

REVENUE STREAM	HOW MONEY COMES IN	PRICE POINT
Consumer SaaS	Monthly subscription (Free / Pro / Business)	Rs 0 – Rs 1,999/mo
B2B Add-on	OmniPilot voice layer on top of BizBot/AdPilot/QuickBillPro	Rs 3,000 – 8,000/mo per product
Enterprise	Custom integrations + team features + dedicated support	Rs 15,000+ /mo + Rs 50K setup
Hardware (future)	OmniPilot wearable pendant	Rs 5,000 – 8,000 one-time

1.8 Top Risks


RISK	SEVERITY	MITIGATION
Local LLM quality insufficient	Medium	Pluggable AI: local default, cloud upgrade available. Quality improving every quarter
"Cool demo, never opened again"	High	Integrate into EXISTING daily workflows (BizBot, invoicing). Don't create new habits
7 integrations = execution complexity	Medium	Start with BizBot only. Add products one at a time. Each must prove value independently
Whisper accuracy in noisy/Hindi	Medium	Use Sarvam AI for Indic languages. Whisper for English. Fallback to text input
Privacy regulation changes	Low	Local-first = naturally compliant. No data to regulate if nothing leaves the device

1.9 What Does Success Look Like in 6 Months?


500 ACTIVE USERS	50 PAYING SUBSCRIBERS	Rs 2.5L MONTHLY REVENUE	3+ PRODUCT INTEGRATIONS	4.5+ APP STORE RATING
----------------------------	---------------------------------	-----------------------------------	-----------------------------------	---------------------------------

Business Design: How It Works

2.1 Revenue Model & Unit Economics



2.2 Cost Structure & Path to Profitability




KEY INSIGHT

Local-first = near-zero infrastructure cost. Unlike competitors who pay \$0.01-0.05 per audio minute for cloud STT, OmniPilot processes everything on-device. The primary cost is the dev team, not servers. Breakeven at ~Month 8 with 100 paying users.

2.3 Customer Journey



2.4 Business Model Canvas



2.5 Channels & Key Partnerships

CHANNEL	HOW IT WORKS	CAC ESTIMATE
BizBot upsell	Existing BizBot customers get OmniPilot offer in-app	Rs 0 (existing relationship)
WhatsApp virality	"Powered by OmniPilot" in auto-replies, referral program	Rs 50-100
App Store organic	SEO: "AI assistant India", "voice assistant business"	Rs 200-500
LinkedIn / Twitter	Founder-led content, product demos, use case stories	Rs 100-300
Partnership: CA/accountant networks	QuickBillPro + OmniPilot bundle for CA client management	Rs 500-1000

Key Partnerships


PARTNER	WHAT WE GET	WHAT THEY GET
WhatsApp (Meta)	Business API access, distribution channel	More businesses using WhatsApp Business
Ollama / whisper.cpp	Free AI models, community support	Enterprise validation of their OSS stack
Sarvam AI	Hindi/Indic language STT models	India distribution, production usage data
Apple (CoreML)	Hardware acceleration for on-device AI	Showcase for Apple Silicon AI capabilities

2.6 Key Metrics (Day 1)


METRIC	TARGET	WHY IT MATTERS
Daily Active Rate	>60%	Must become a daily habit, not a toy
Queries per user per day	>5	Proves value in workflow, not just novelty
Memory recall accuracy	>85%	"What did I discuss?" must return useful results
Action completion rate	>90%	"Send invoice" must actually send it, not fail silently
Month-3 retention	>70%	Long-term stickiness, not just trial curiosity

Product Design: What We Build

3.1 System Architecture




3.2 Data Flow Pipeline




PRIVACY PROMISE

End-to-end latency under 1 second. All processing happens on your device. No audio, text, or memory data ever touches a cloud server. Your conversations are yours alone.

3.3 Memory Architecture



3.4 Tech Stack & RAM Budget



Component	Tool	RAM	Latency
Voice Detection	Silero VAD (ONNX)	50MB	<1ms per 30ms chunk
Speech-to-Text	whisper.cpp + CoreML	1.5GB	~200ms per 3-5s chunk
AI Brain	Ollama + Qwen3 8B Q4	5GB	20-30 tokens/sec
Memory + Search	SQLite + sqllite-vec + FTS5	100MB	<50ms per query
Embeddings	all-MiniLM-L6-v2 (ONNX)	80MB	<10ms per embedding
Wake Word	Picovoice Porcupine	10MB	<1ms


Total RAM: ~6.7GB on M1 16GB Mac, leaving ~5GB free headroom + OS

3.5 MVP Feature Set (v1.0)

Feature	In MVP	Post-MVP
Always-on mic with VAD	Yes	
Push-to-talk hotkey	Yes	
Local transcription (Whisper)	Yes	
Memory storage + search	Yes	
"What did I discuss about X?"	Yes	
Daily end-of-day summary	Yes	
WhatsApp voice note processing		Phase 2
BizBot voice commands		Phase 4

FEATURE	IN MVP	POST-MVP
Flutter mobile app		Phase 3
Cross-product commands		Phase 5
Proactive reminders		Phase 5
Hardware pendant		Phase 6

3.6 Product Portfolio Integration



PRODUCT	SYNERGY	TOP USE CASE	REVENUE ADD-ON
BizBot	95%	Voice-to-WhatsApp: "Reply to that customer"	Rs 5-8K/mo
DevPilot	85%	Voice sprint management: "Mark task done"	Internal tool
QuickBillPro	80%	Voice invoicing: "Invoice Brand Tusker Rs 2.3L"	Rs 3-5K/mo
AdPilot	75%	Voice metrics: "What's my campaign ROAS?"	Rs 3-5K/mo
TradePilot	70%	Ambient trading: "How's my portfolio?"	Rs 3K/mo
SetlIn	65%	City settling: "Find an electrician nearby"	Rs 2K/mo
DesignPilot	55%	Voice design: "Modern kitchen 10x12"	Rs 2K/mo

CROSS-PRODUCT MAGIC

"Brand Tusker's ad campaign is doing well. Send a thank-you on WhatsApp and draft next month's invoice."

One voice command triggers **AdPilot** (pull metrics) + **BizBot** (send WhatsApp) + **QuickBillPro** (generate invoice). No competitor can do this.

Execution: Getting It Done

4.1 Build Roadmap



4.2 Sprint 1: Desktop Companion (Week 1-2)

DAY	TASK	DELIVERABLE
Day 1	Install Ollama + Whisper.cpp + CoreML backend	Both running locally, tested
Day 2	Swift menu bar app skeleton + CoreAudio mic capture	App captures audio, sits in system tray
Day 3	Silero VAD integration + hotkey (push-to-talk)	App detects speech, ignores silence
Day 4-5	Whisper transcription pipeline + SQLite memory schema	Spoken words become searchable text in DB
Day 6-7	sqlite-vec embeddings + FTS5 hybrid search	"What did I discuss about X?" works
Day 8-9	Ollama integration: summarize, extract entities/people	Auto-tags people, topics, decisions
Day 10	End-of-day summary generation	Daily digest of everything discussed
Day 11-12	Polish: wake word, settings UI, memory browser	User can browse/search all memories
Day 13-14	Testing, bug fixes, dogfooding	v0.1 ready for daily personal use

4.3 Launch Checklist

CATEGORY	ITEM	STATUS
Legal	Privacy policy (local-first, no data collection)	Required for app store
Legal	Terms of service	Required for app store
Legal	WhatsApp Business API approval	Phase 2 (apply early)
Tech	Ollama + Whisper.cpp bundled installer	One-click Mac setup
Tech	Auto-update mechanism	Sparkle framework (macOS)
Marketing	Landing page at omnipilot.in	Before any public launch
Marketing	Demo video (60s)	Shows voice query → answer flow
Analytics	Usage tracking (local, privacy-safe)	Query count, daily active, features used
Support	Discord community or WhatsApp group	Early adopter feedback channel

4.4 Go-to-Market Plan

WEEK	ACTION	CHANNEL	GOAL
Week 1-2	Build & dogfood	Internal	v0.1 working for personal use
Week 3-4	Private alpha (10 people)	Friends + Sidewall network	First external feedback
Month 2	BizBot customer beta	In-app notification	50 BizBot users try OmniPilot
Month 3	Public launch	Product Hunt + LinkedIn + App Store	500 downloads, 50 paid
Month 4-6	Growth loops	WhatsApp virality + referral program	2000 users, 200 paid

4.5 Risk Mitigation Playbook

RISK (FROM PHASE 1)	BACKUP PLAN	TRIGGER TO ACTIVATE
Local LLM quality insufficient	Pluggable AI: add Claude API as premium tier	Memory recall accuracy <70%
"Cool demo, never opened again"	Daily summary push notification + BizBot integration hooks	DAU drops below 30%
7 integrations = execution overload	Focus on BizBot only for 3 months. Others wait	Sprint velocity drops

RISK (FROM PHASE 1)	BACKUP PLAN	TRIGGER TO ACTIVATE
Hindi/Indic accuracy poor	Partner with Sarvam AI. Offer English-first, add languages	Hindi WER >20%
Apple rejects menu bar app	Distribute via website (DMG). Skip App Store initially	App Store rejection

READY TO BUILD

Phase 1 starts tomorrow (April 17, 2026). Day 1: Install Ollama + Whisper.cpp. Day 2: Swift menu bar skeleton. By Day 14: a working personal AI companion on your Mac that listens, remembers, and answers.

OmniPilot

Your voice. Your data. Your products.
One AI to manage everything.

Strategy ✓

Business ✓

Product ✓

Execution ✓



Sidewall / DevPilot

soumya@sidewall.in | support@devpilot.co.in